

When AI Joins the Org Chart

Cycle: 12 May 2026 · Audience: Senior strategy, innovation, public-sector foresight and capital-allocation professionals · Horizon: 12 to 18 months

Read estimate: 2 min Executive Synthesis · 22 min full read

Executive Synthesis

Three years after generative AI moved from demo to workflow, the strategic question is no longer whether to deploy AI. It is which seat AI occupies on the organisation chart: a tool beside a human, an autonomous worker reporting to a manager, or a workflow whose human owner has gone unclaimed. This cycle establishes the baseline, and the baseline is a chart being redrawn live while the infrastructure that should support it, trust, talent and capital, struggles to keep pace.

Four numbers landing within five months of each other describe the shift. Per the [McKinsey State of AI 2025](#) (November 2025), 23% of organisations are now scaling at least one agentic AI system, yet only 39% report any enterprise-level EBIT impact and only 6% qualify as AI high performers. Per the [PwC 29th Global CEO Survey](#) (January 2026), just 12% of CEOs say AI has delivered both cost and revenue benefits, 56% report no significant financial benefit, and the gap between the AI haves and have-nots is now a top CEO concern. Per the [Stanford AI Index 2026](#), employment for software developers aged 22 to 25 has fallen nearly 20% from 2024, in the same domains where productivity gains are appearing. And per [Goldman Sachs Global Institute](#) (May 2026), the build-out now sits on a \$7.6 trillion six-year capex envelope whose supply-side assumptions are more sensitive than demand.

*The strategic question of 2026 is no longer whether to deploy AI.
It is which seat in the org chart AI occupies, and which seats it
vacates.*

Read together, these are not separate headlines. They describe one structural change with two halves: AI is moving from tool to colleague to manager-of-tasks, and the supporting infrastructure is lagging it. The spine of this cycle is the deployment shift from

copilots to agents, because it is the change that forces every other question. Once an agent holds a permission set and runs a multi-step workflow, an organisation needs an owner for it, a verification stack to trust it, a talent pipeline that still produces judgement when AI does the routine work, and a capital plan that knows what the spend is supposed to return. The surprise for many leadership teams is how far the deployment has run ahead of the accountability: organisations are hiring agents the way they once hired outsourcing providers, but without the equivalent of the BPO contract.

Where this analysis could be wrong

The central assumption is that agentic AI is becoming a workforce category, not just a better productivity tool. If, twelve months from now, Klarna-style reversals dominate enterprise headlines and EBIT impact stays stuck near McKinsey's 6% high-performer share, the thesis is wrong: AI is a tool with limits and the org-chart framing is premature. The disconfirming evidence would be concrete and visible, a run of high-profile agent rollbacks plus a second flat year of enterprise EBIT data. If instead agent deployment crosses 40% of enterprise applications, EU AI Act enforcement lands in the autumn, and entry-level hiring stays depressed for a second year, the thesis hardens into a structural feature of the operating model.

The decisions that cannot be deferred this cycle

1. **Decide whether you have an agent operating model, not just an AI strategy.** Licences are not deployment; named owners, permission sets and escalation contracts are.
2. **Move verification from policy artefact to procurement capability.** The EU AI Act becomes fully applicable on 2 August 2026, and trust is now something you buy and audit, not something you write down.
3. **Take a deliberate position on the entry-level rung.** AI is hollowing out the training ground; quietly cutting graduate hiring and deliberately redesigning it look identical for a year and very different at three.

Each is developed below, with a decision posture, in the four Strategic Implications.

Audience Snapshots

Four lenses on the same intelligence base, one per buyer-team that most often asks Shaping Tomorrow for AI foresight. Each card surfaces the one question this cycle puts to that function.

OPERATE · BOARDS, CEOS, RISK, STRATEGY

Is our AI investment translating into EBIT, or absorbing budget while it learns?

The shift: only 12% of CEOs say AI has delivered both cost and revenue benefits and 56% report no significant financial benefit, while CEOs with Responsible AI frameworks and enterprise-wide integration are roughly three times more likely to report meaningful returns.

The question to brief: is the AI line on next year's plan backed by a workflow-redesign sub-line, a named model-risk owner and a defined EBIT measurement method, or are we funding the experiment and calling it the return?

BUILD · INNOVATION, TECHNOLOGY, R&D STRATEGY

Are we building agent infrastructure for 2027, or copilot integrations commoditised by year-end?

The shift: 23% of organisations are scaling at least one agentic system and Gartner projects up to 40% of enterprise applications will feature task-specific agents by the end of 2026, making multi-agent orchestration the next platform-architecture decision.

The question to brief: are identity, permissioning, action-logging, containment and rollback being treated as product requirements now, or still as security afterthoughts we will retrofit later?

GOVERN · PUBLIC-SECTOR FORESIGHT, RESILIENCE

Is our AI deployment moving from pilot to integration, or stuck where the data foundations failed?

The shift: the UK has scaled its Minute AI scribe to 1,000 officials across 22 local authorities, but the pilot-to-integration gap remains the dominant public-sector constraint, and the EU AI Act becomes fully applicable on 2 August 2026.

The question to brief: does our AI risk register treat both the model-behaviour problem and the workforce problem, or are we still scoping AI as an IT procurement?

ALLOCATE · INVESTORS, ASSET MANAGERS, CAPITAL ALLOCATION

At \$7.6 trillion of cumulative AI capex to 2031, what assumption is the portfolio actually betting on?

The shift: Goldman Sachs projects \$765bn of AI capex in 2026 rising to \$1.6tn by 2031, with hyperscaler free cash flow already compressing, and the four supply-side assumptions more sensitive than demand to whether the bill is met by returns.

The question to brief: have we built a portfolio-level depreciation-schedule sensitivity and a hyperscaler free-cash-flow stress test, or is "AI exposure" still a stock-picking label?

Signal Clusters


The cycle's signals are organised into five clusters, ranked by impact on the decisions leadership teams face now. **Immediate**: changes the operating model, the corporate structure or the proposition. **Near-Term**: changes competitive position over the next twelve months. **Longer-Range**: a multi-year structural factor to track and revisit each cycle.

1. The agent deployment shift

IMMEDIATE

AI has stopped being a feature inside someone else's software and started being the software. In 2025, "scaling AI" meant rolling out Copilot licences; in 2026 it means deploying autonomous agents that run multi-step workflows, hold permissions and report to human owners. The deployment looks less like SaaS adoption and more like outsourcing, and that is the point: organisations are hiring agents the way they used to hire BPO providers, which means the 2026 capability question is not "do we have an AI strategy" but "do we have an agent operating model".

Adoption is broad; agent deployment is narrow



The deployments that are in place work hard; most of the enterprise still runs on the copilot model.

Sources: McKinsey State of AI 2025; Stanford AI Index 2026.

- AI use in at least one business function has reached 88% of organisations, up from 78% a year earlier, but only 23% are scaling an agentic system and 39% are experimenting. [McKinsey State of AI 2025](#) (November 2025).
- Agent deployment remains in the single digits across nearly all business functions even as organisational adoption reaches 88%: the deployments in place are concentrated and working hard, in IT, knowledge management, contact centres and software. [Stanford AI Index 2026](#) (April 2026).
- The Klarna arc is the case study every operator should have read: a 2024 deployment claimed the work of 700 agents, but by mid-2025 Klarna was re-hiring humans and positioning human service as a premium tier above the AI layer. The lesson is not that AI fails in service; it is that the deployment model has to specify which interactions belong to which class of worker. [OpenAI / Klarna](#) (2024); [Customer Experience Dive](#) (2025).

- The hyperscalers are institutionalising agent deployment as a category, with governance, identity and audit primitives built into the platform: NVIDIA and ServiceNow's autonomous-agent partnership and Google's Gemini Enterprise Agent Platform, with Tata Steel building a 300-agent fleet. [NVIDIA and ServiceNow](#) (May 2026); [Google Cloud](#) (April 2026).

Counter-argument

The agent story may be a vendor-led narrative that overstates deployment maturity. The Klarna reversal is one of several, model reliability is not yet production-grade for high-stakes workflows, and the UK AI Security Institute notes universal jailbreaks have been found in every frontier model tested. Agentic AI may turn out to be the next "self-driving in two years", perpetually nearly ready, delivering real value in narrow, well-instrumented workflows for years before it becomes the dominant pattern.

Decision link: Strategic Implication [1](#).

2. The verification stack becomes a procurement requirement

IMMEDIATE

Trust in AI outputs is no longer a behavioural posture. It is becoming an infrastructure stack: identity for agents, logging for actions, classifiers for output, evaluations for capability and external review for accountability, with regulatory enforcement landing in August 2026. For operators in the EU and EU-adjacent markets, that is a hard deadline; for everyone, the trust stack has moved from a policy item to a procurement item, something an organisation buys, audits and is expected to read.

- The EU AI Act becomes fully applicable on 2 August 2026, with high-risk system rules, transparency rules and Commission enforcement for general-purpose AI models all entering force the same month. [European Commission](#) (January 2026).
- The standards layer is consolidating around the same shape: the OECD's Due Diligence Guidance for Responsible AI gives multinationals a structured playbook, and NIST has announced an AI Agent Standards Initiative naming agent identity, action logging and containment boundaries as the three pillars of agent governance. [OECD](#) (February 2026); [NIST](#) (2026).
- Capability and safeguard are moving together, not apart: the UK AI Security Institute reports agentic autonomy roughly doubling every two months, universal jailbreaks found in every system tested, and AI models now completing apprentice-level cyber tasks around 50% of the time, up from roughly 10% in early 2024. [UK AI Security Institute](#) (2026).
- The disclosure default is converging on capability evaluation plus safety evaluation plus residual risk plus mitigation: Anthropic's Responsible Scaling Policy v3.0 commits to public Risk Reports every three to six months and external review at higher capability levels, a

pattern operators should expect to receive and to read the way they read a SOC 2 audit. [Anthropic](#) (February 2026).

Counter-argument

Verification infrastructure may be moving fast enough to look impressive but slow enough to be irrelevant to the operational risk. The hallucination problem is unsolved at the model layer, agent containment is unsolved at the system layer, and shadow AI use, employees pasting confidential data into personal chatbots, is unaddressed at the data-governance layer. A trust stack that does not cover the day-to-day failure modes is a compliance artefact, not a risk control, and boards that treat AI Act compliance as the ceiling rather than the floor will find the gap when something embarrassing reaches the front page.

Decision link: Strategic Implication 2.

3. The payoff gap: AI's haves and have-nots

NEAR-TERM

The deployment is real, but the return is not yet evenly arriving. Only a small share of organisations can show enterprise-level financial impact from AI, and the gap between the firms that can and the firms that cannot has become the dominant strategic anxiety of the cycle. This is a competitive-positioning question over the next twelve months: the divergence is widening while most management teams cannot yet name the use cases AI is supposed to be paying for.

- Just 12% of CEOs say AI has delivered both cost and revenue benefits and 56% report no significant financial benefit, while the gap between the AI haves and have-nots is now a top CEO concern, ahead of innovation or viability. [PwC 29th Global CEO Survey](#) (January 2026).
- Only 39% of organisations report any enterprise-level EBIT impact from AI and only 6% qualify as AI high performers, a thin leading edge against broad adoption. [McKinsey State of AI 2025](#) (November 2025).
- Goldman Sachs finds no meaningful relationship between AI adoption and productivity at the economy-wide level, but a median 30% gain in two narrow use cases, with only 10% of S&P 500 management teams quantifying AI's impact on specific use cases and only 1% quantifying earnings impact: the measurement gap is the real gap. [Fortune, reporting Goldman Sachs](#) (March 2026).
- UK firm-level data shows large firms at 36 to 44% AI adoption and 56% of AI-using firms reporting self-assessed productivity gains, but 95% of AI-using SMEs report no change in workforce size: adoption is running well ahead of measured structural effect. [Bennett School, Cambridge](#) (2026).

Counter-argument

The payoff gap may be a measurement problem, not a returns problem. The productivity J-curve, the pattern where measured productivity dips before gains from a general-purpose technology show up in national accounts, is well established for electricity and computing. The "no enterprise-wide EBIT impact yet" data looks much like the early data on enterprise software in the early 1990s. Acemoglu's conservative macro ceiling sits near +0.07 percentage points of annual total factor productivity, but if the J-curve plays out as before, the 2025 to 2027 capex funds the productivity surge of 2028 to 2032, and the "reckoning" reads as impatience rather than a structural problem.

Decision link: Strategic Implication 4.

4. The entry-level reset

NEAR-TERM

AI is not displacing the workforce evenly. It is hollowing out the entry-level rung of the talent pyramid, the rung organisations have always used to renew themselves, because the tasks that used to be how junior staff learned (cleaning data, drafting summaries, first-pass code, triage) are precisely the tasks AI does well. The structural story is not "AI replaces workers"; it is "AI replaces the training ground", and the hiring decisions that respond to it are being made over the next twelve months.

- Employment for software developers aged 22 to 25 has fallen nearly 20% from 2024, concentrated in the cohort whose first jobs are the tasks AI does best, while US entry-level postings are down around 35% over 18 months. [Stanford AI Index 2026](#) (April 2026); [WEF](#) (March 2026).
- The work assumed to be done by AI in early-career roles is being pushed upward, leaving middle management and senior talent overextended: the pyramid is not flattening, it is hollowing in the middle and stretching at the top. [Yale Insights](#) (2026).
- The entry-level jobs that are being filled increasingly require AI fluency on day one: 4.2% of full-time early-career postings now mention AI skills, nearly double the share a year earlier. [CNBC and Indeed Hiring Lab](#) (April 2026).
- The counter-pattern exists: a field experiment with around 5,200 customer-support agents shows AI assistance raising issues resolved per hour by 15% on average and up to 35% for the least-experienced workers, evidence that AI can upskill entry-level staff when the deployment is configured for assistance rather than replacement. [Brynjolfsson, Li and Raymond](#) (2023).

Counter-argument

The entry-level squeeze may be cyclical, a post-pandemic correction in over-hired technology functions rather than a permanent AI-driven structural shift. Healthcare entry-level postings rose over the same period that software entry-level postings fell, so the sectoral mix matters, and NACE's spring update revised the Class of 2026 forecast upward as employers re-balanced toward using AI to lift entry-level productivity rather than replace it. If the tech hiring cycle turns in 2027, the "AI is destroying entry-level jobs" narrative may prove a measurement artefact of where AI was deployed first and hardest.


Decision link: Strategic Implication 3.

5. The capex reckoning

LONGER-RANGE

The AI capex bill is now visible and large, and hyperscaler free cash flow is being compressed in real time while the productivity the spend is meant to fund is still confined to narrow use cases. This is a multi-year structural factor rather than a this-quarter event, but it sets the backdrop for every organisation with material AI capex and every investor pricing AI exposure: the question is no longer whether AI demand materialises, it is whether the supply-side assumptions hold and the bill is met by returns.

The six-year AI capex envelope



Goldman Sachs Global Institute baseline. The 2028 bar is indicative of the rising path between the 2026 and 2031 anchors.

- The cumulative six-year AI capex envelope sits at \$7.6 trillion on the baseline, \$765bn in 2026 rising to \$1.6tn by 2031, with silicon useful life identified as the single most influential variable: shorten effective useful life from six years to four and cumulative spend rises by hundreds of billions. [Goldman Sachs Global Institute](#) (May 2026).
- The four hyperscalers will spend a combined roughly \$700bn on AI infrastructure in 2026 alone. [Fortune](#) (April 2026); [Yahoo Finance and S&P Capital IQ](#) (May 2026).

- Free cash flow is turning negative under the weight of the build-out: Amazon's 2026 FCF deficit is projected at \$17bn to \$28bn and Meta's FCF is expected to fall around 90%, even as Google Cloud Q1 revenue rose 63% year on year and Microsoft's AI business runs at an annualised \$37bn. [CNBC](#) (February 2026).
- The macro literature offers opposing anchors: Acemoglu sets a conservative ceiling near +0.07 percentage points of annual total factor productivity, while the IMF's scenario work warns the gap between AI-prepared and AI-unprepared economies is widening, with fiscal, monetary and financial-stability consequences. [Acemoglu](#) (2024); [IMF](#) (April 2026).

Weak signals to watch

- **WEAK SIGNAL** Sovereign credit and AI exposure: the IMF's scenario note and the BIS's work on AI's financial-stability channels both point toward something rating agencies have not yet formalised, that jurisdictions and corporates with heavy AI capex commitments and weak revenue-conversion stories may begin to receive different treatment in credit spreads. Would gain weight on the first sovereign rating action or major corporate downgrade citing unmonetised AI capex.
- **WEAK SIGNAL** An external-review market for AI: Anthropic's commitment to appoint expert third-party reviewers with unredacted access to Risk Reports is the seed of an audit-style review market. Would gain weight if the EU AI Act's market-surveillance architecture pulls a small group of credible external reviewers into being over the next 18 to 24 months.

Counter-argument

The capex may be more disciplined than the free-cash-flow headlines suggest. It is being underwritten by real, fast-growing cloud revenue and a large order backlog, and the FCF compression is a deliberate investment choice by some of the most cash-generative companies in the world, not a sign of distress. Google Cloud's 63% revenue growth and Microsoft's \$37bn annualised AI run-rate are demand signals, not hopes. On this reading the "reckoning" is a sell-side framing, and the build-out is a rational land-grab in a market whose demand curve is still bending upward.

Decision link: Strategic Implication [4](#).

Strategic Implications

Four decisions turn this cycle's signals into next year's plan. Each names the move, a horizon and a decision posture. Opportunities are expressed here, as Prepare-posture moves, rather than in a separate section.

SI 1: Build the agent operating model before the agent fleet

The leadership team should, this planning cycle, decide and document the agent operating model rather than expand the licence count, because licences are not deployment and the organisations scaling agents are encoding ownership, permissions and escalation into the platform itself. Three artefacts make the model real: an agent inventory with named owners, permission sets and escalation rules; a "human-managed agent" job description, because the agent owner is now a real role; and a model-risk register that integrates AI Act high-risk classifications and NIST profiles. An organisation that cannot answer who owns, what is permitted and where the escalation contract sits for its top five deployed agents is still pre-deployment, regardless of how many licences it holds.

Action: by the next planning round, an agent inventory with named owners and escalation contracts for the top five deployed agents, signed off by an accountable executive.

DECIDE

Draws on Cluster 1.

SI 2: Treat verification as a procurement capability, not a policy artefact

The leadership team should, ahead of the EU AI Act's 2 August 2026 applicability date, move verification from a written policy into the procurement function, because the trust stack is now buyable and auditable and August 2026 is a hard deadline for EU and EU-adjacent operators. Concretely: require Risk Reports, model cards and incident-disclosure terms as standard procurement clauses; map AI-touching workflows against the high-risk classifications; and identify which three would not pass an external review today. The OECD Due Diligence Guidance and the EU GPAI Code of Practice give the template even for non-EU operators. The upgrade is low-cost and high-leverage, and treating compliance as the floor rather than the ceiling is the difference between a risk control and a compliance artefact.

Action: by end-Q3 2026, AI verification clauses embedded in standard procurement, and a named closing plan for the three highest-exposure workflows.

DECIDE

Draws on Cluster 2.

SI 3: Redesign the entry-level for AI-assisted apprenticeship, not AI-replaced production

The leadership team should, in the workforce-planning cycle, take a deliberate position on the entry-level rung rather than let graduate hiring drift down by default, because AI is replacing the training ground and the classical apprenticeship model has no obvious replacement. Three moves follow: redesign first-year roles around AI-output review, exception handling and judgement tasks rather than routine production; make AI fluency a hiring screen rather than an optional skill; and explicitly cost the senior-team burnout that comes from absorbing junior tasks into the AI business case. The counter-pattern, exemplified by firms that are hiring more graduates because juniors ramp fastest on AI tools, is rational but not yet widely copied, and the cost of skipping the third move shows up later as turnover.

Action: by the next workforce-planning cycle, a redesigned first-year role profile and an explicit, costed position on graduate intake.

PREPARE

Draws on Cluster 4.

SI 4: Build the AI-return stress test into the plan and the investment thesis

The leadership team, and investors holding AI-exposed portfolios, should build a standing stress test of whether the AI spend is earning its return, because the payoff gap is now the dominant strategic concern and the capex envelope is large enough that the supply-side assumptions matter more than demand. For operators: name the specific use cases quantified for cost or revenue impact, and state the exposure to a twelve-month delay in AI revenue materialising. For investors: build a portfolio-level depreciation-schedule sensitivity around silicon useful life and a hyperscaler free-cash-flow recovery stress test. The work is mechanical and the inputs are knowable line items; the discipline is doing it before the market reprices the divergence rather than after.

Action: by year-end 2026, a quantified AI-return stress test, use-case by use-case for operators and depreciation-and-FCF sensitivity for investors.

PREPARE

Draws on Clusters 3 and 5.

Scenario Matrix

Two uncertainties will shape the next eighteen months more than any other. First, the pace of agent deployment: does the move from copilots to production agents accelerate, or stall on reliability and governance friction? Second, the productivity payoff: does enterprise-level EBIT impact materialise, or does the capex outrun the returns? The four scenarios below are planning aids, not forecasts.

The Working Org Chart

High agents, high payoff. Agent deployment crosses 40% of enterprise applications, EU AI Act enforcement creates a verification floor that becomes a procurement standard, and the productivity gains visible in support, software and marketing spread to finance and HR. Enterprise EBIT impact moves from a 6% high-performer share toward a quarter of organisations, hyperscaler capex is underwritten by cloud revenue growth, and the talent pipeline resets around AI-fluent juniors who arrive ready to manage agents.

Early indicators: agent deployment scaling beyond IT into customer-facing and finance functions; credible cross-firm productivity studies showing sectoral gains; hyperscaler free cash flow normalising in the second half of 2026.

Pilot Purgatory

Low agents, low payoff. Agent deployments stall on reliability, hallucination and governance friction, Klarna-style reversals become the dominant pattern in customer-facing functions, and the "no significant benefit" majority persists into 2027. Hyperscaler capex continues but free-cash-flow disappointment triggers the first downgrades, the EU AI Act is read as a compliance ceiling rather than a risk floor, and the entry-level squeeze hardens without a replacement training architecture emerging.

Early indicators: multiple high-profile agent rollbacks; hyperscaler capex cuts in 2026 earnings calls; rating-agency commentary on AI-capex sustainability.

Bifurcated Enterprise

High agents, uneven payoff. A small group of AI high performers widens from 6% toward 15 to 20% and pulls away on margin, market share and innovation cadence, while the laggards' returns disappoint. The bifurcation is sectoral, with technology, financial services and professional services pulling ahead while manufacturing, the public sector and the mid-market lag. Talent flows accelerate the gap, EU AI Act enforcement bites unevenly, and capital markets begin to price the divergence directly.

Early indicators: widening valuation gaps between AI-deployed and AI-undeployed peers; M&A consolidation driven by AI-capability arbitrage; CEO surveys showing a sharper bimodal split on AI returns.

Regulated Slowdown

Low agents, high payoff in narrow segments. EU AI Act enforcement and analogous US state-law enforcement create real procurement friction, agent deployment slows by twelve to eighteen months in regulated industries, and the productivity gains that do materialise concentrate in unregulated knowledge work such as software and marketing. Public-sector AI deployment slows further, and the narrative shifts from "AI everywhere" to "AI where it does not need a regulator's permission".

Early indicators: a first major EU AI Act enforcement action in the second half of 2026; regulated-industry AI deployment data flatlining; "AI exception" carve-outs appearing in major procurement contracts.

What We Are Not Planning For

A frontier-model capability discontinuity

The base case is that capability improvements are continuous and measurable. A discontinuous step, a model markedly more capable than the current frontier on multi-step reasoning, autonomous research or persuasion, would re-anchor every cluster in this report. The probability is hard to bound, but the consequence would be high enough that it is tracked rather than planned around.

Reinstate if: the UK AI Security Institute or a credible third party reports a step-function on benchmark evaluations, or a non-public model achieves expert-level autonomy on a strategically significant task.

A regulatory rollback in the US triggering bifurcated global compliance

The base case is that the EU and US regimes diverge gradually rather than splitting hard. A genuine US rollback, federal pre-emption of state AI laws or withdrawal from international AI-safety institutions, would force operators to carry two compliance stacks and would change the procurement calculus in cluster 2. It is a watched risk, not a planning assumption.

Reinstate if: federal pre-emption legislation passes, or the US withdraws from a major international AI-safety agreement.

A hyperscaler capex retrenchment ahead of cloud-revenue maturity

The base case is that hyperscaler capex stays funded by cloud growth and backlog through 2026. A sharp retrenchment, a hyperscaler cutting projected spend by 20% or more in a single guidance cycle, would invalidate the cluster 5 framing as written and accelerate the Pilot Purgatory scenario. The cluster tracks the compression; this exclusion tracks the tail.

Reinstate if: any one of the four hyperscalers cuts AI capex guidance by 20% or more in a 2026 reporting cycle.

Discussion Points for the Leadership Team

1. Which seat does AI occupy on our org chart today, and which seats do we expect it to occupy in eighteen months? If the honest answer is "a Copilot rollout and an AI policy", it is too early to call that an AI strategy.
2. Of our top five deployed AI use cases, which are quantified for cost or revenue impact by name, and what is the plan to close that measurement gap by year-end rather than expand the pilot list?
3. Where in our talent pipeline are we absorbing AI's effect on entry-level work: are we quietly cutting graduate hiring, or deliberately redesigning first-year roles? The two paths look identical for eighteen months and very different at thirty-six.
4. Is our AI risk register integrated with the EU AI Act high-risk classifications and the NIST framework, or are they parallel documents nobody cross-references?
5. What is our one-line answer to "if AI capex does not deliver economy-wide productivity by 2028, what changes in our plan", and what is our depreciation-schedule and free-cash-flow sensitivity if it does not?

Source Confidence Register

This cycle drew on 39 verified sources across the five clusters, each checked by WebSearch and WebFetch during compilation, on a soft three-month recency window with named structural anchors permitted (Acemoglu 2024, Brynjolfsson and colleagues 2023, OECD Employment Outlook 2023, IMF Working Paper 2025/067). The tier mix is disclosed honestly: clusters 1 and 2 are anchored on Tier 1 regulators and standards bodies and Tier 2 institutional research, with Tier 4 AI-lab and vendor sources used only for supply-side and disclosure signals, never as the sole anchor of a board-level claim. Cluster 5 leans more heavily on Tier 3 quality journalism, because the hyperscaler capex and free-cash-flow picture is reported through earnings journalism aggregating primary corporate disclosures. The Klarna case is held with both its original vendor source and its corrective trade-press source so the operational lesson can be drawn cleanly, and the macro debate is presented as live, with the conservative Acemoglu bound held against the consensus build-out frame. Tier 1 and Tier 2 sources together are 26 of 39, around 67%.

Source tiers: **Tier 1**, governments, regulators and intergovernmental bodies. **Tier 2**, think-tanks, academic institutes, major consultancies and quality data providers. **Tier 3**, quality journalism and specialist trade press. **Tier 4**, vendor, company and practitioner sources, used only as directional corroboration.

Cluster 1: The agent deployment shift

Source	Tier	Date	Key claim used
McKinsey, The State of AI 2025	Tier 2	Nov 2025	88% of organisations use AI in at least one function (up from 78%); 23% scaling an agentic system, 39% experimenting; 6% are AI high performers.
Stanford HAI, AI Index 2026 (Economy)	Tier 2	Apr 2026	Agent deployment in single digits across nearly all functions despite 88% adoption; productivity gains concentrated in support, software and marketing.
WEF, Where is AI moving beyond experimentation	Tier 2	Mar 2026	Field reports from six enterprise leaders on which deployments are moving from pilot to production.
Forrester, AI Agents Predictions (Oct 2025)	Tier 2	Oct 2025	2026 as the breakthrough year for multi-agent systems; enterprise software shifting to agent-first design.
Gartner, 40% of enterprise apps to feature AI agents	Tier 4	Aug 2025	Up to 40% of enterprise applications to feature task-specific AI agents by end-2026, from under 5% in 2025.

Source	Tier	Date	Key claim used
OpenAI and Klarna case study	Tier 4	2024	Klarna AI assistant doing the equivalent work of 700 agents; the original deployment claim in the Klarna arc.
Customer Experience Dive, Klarna changes its AI tune	Tier 3	2025	Klarna re-hiring humans; human service repositioned as a premium tier above the AI layer.
NVIDIA and ServiceNow, autonomous AI agents	Tier 4	May 2026	Project Arc announced at ServiceNow Knowledge 2026; vendor alliance institutionalising a production agent platform with governance primitives.
Google Cloud, Gemini Enterprise Agent Platform	Tier 4	2026	Enterprise agent platform with governance and audit primitives; Tata Steel building a 300-agent fleet.

Cluster 2: The verification stack becomes a procurement requirement

Source	Tier	Date	Key claim used
European Commission, AI Act overview and timeline	Tier 1	Jan 2026	AI Act fully applicable 2 August 2026; high-risk, transparency and GPAI enforcement entering force that month.
European Commission, GPAI Code of Practice	Tier 1	Jul 2025	Voluntary compliance tool for general-purpose AI providers; a template usable by non-EU operators.
OECD, Due Diligence Guidance for Responsible AI	Tier 1	Feb 2026	Structured AI risk-management playbook for multinationals.
NIST, AI Risk Management Framework	Tier 1	2026	US baseline for AI risk governance; AI Agent Standards Initiative naming identity, action logging and containment as the three governance pillars.
UK AI Security Institute, Frontier AI Trends Report	Tier 1	2026	Agentic autonomy doubling roughly every two months; universal jailbreaks in every system tested; apprentice-level cyber task success around 50%, up from roughly 10% in early 2024.
BIS, financial-stability implications of AI	Tier 1	Jan 2026	Three financial-stability channels for AI: market functioning, operational concentration, propagation of stress.

Source	Tier	Date	Key claim used
UK Government, AI Opportunities Action Plan: One Year On	Tier 1	Jan 2026	Minute AI scribe scaled to 1,000 officials across 22 local authorities; public-sector pilot-to-integration gap as the dominant constraint.
Anthropic, Responsible Scaling Policy v3.0	Tier 4	Feb 2026	Public Risk Reports every three to six months; external review at higher capability levels; the disclosure-default signal.
OpenAI and Anthropic joint safety evaluation	Tier 4	2025	Cross-lab evaluation precedent likely to inform regulatory practice.
Aon, AI Risk 2026: Practical Agenda	Tier 2	2026	Insurance and risk perspective on AI failure modes: hallucination, data leakage, agent misbehaviour, governance gaps.
Suprmind, AI hallucination statistics and benchmarks	Tier 4	April 2026	362 documented AI incidents in 2025 per the AI Incident Database, up 55% year on year. Vendor aggregator citing third-party benchmarks; used directionally.

Cluster 3: The payoff gap: AI's haves and have-nots

Source	Tier	Date	Key claim used
PwC, 29th Global CEO Survey	Tier 2	Jan 2026	4,454 CEOs: 12% report both cost and revenue benefits, 56% no significant benefit; the AI haves and have-nots gap is the top CEO concern; strong-foundation CEOs around 3x more likely to report returns.
Fortune, reporting Goldman Sachs on the productivity link	Tier 3	Mar 2026	No meaningful economy-wide AI-productivity link; median 30% gain in two narrow use cases; 10% of S&P 500 teams quantify use-case impact, 1% quantify earnings impact.
BlackRock Investment Institute, Q2 2026 Outlook	Tier 2	2026	AI capex contributing around three times the historical investment average to US growth; firms turning AI into durable cost advantage will separate from those that cannot.

Source	Tier	Date	Key claim used
Stanford HAI, AI Index 2026 (Economy)	Tier 2	Apr 2026	Measured productivity gains of 14 to 15% in customer support, 26% in software and 73% in marketing; private AI investment up 127.5% in 2025.
Bennett School, Cambridge, UK firm-level AI adoption	Tier 2	2026	Large UK firms at 36 to 44% AI adoption, micro-businesses at 14%; 56% of AI-using firms report productivity gains, 95% of SMEs report no workforce change.
Acemoglu, The Simple Macroeconomics of AI	Tier 2	2024 (anchor)	Conservative macro ceiling near +0.07 percentage points of annual total factor productivity over a decade.

Cluster 4: The entry-level reset

Source	Tier	Date	Key claim used
Stanford HAI, AI Index 2026 (Economy)	Tier 2	Apr 2026	Employment for software developers aged 22 to 25 down nearly 20% from 2024.
WEF and Cognizant CPO, how AI is changing entry-level work	Tier 2	Mar 2026	US entry-level postings down around 35% in 18 months; Cognizant hired 25,000 graduates in 2025; junior tasks pushed upward into senior workloads.
Yale Insights, job destruction before careers start	Tier 2	2026	AI-driven job destruction concentrated in entry-level and apprenticeship roles, reshaping how the talent pyramid renews.
WEF, Future of Jobs Report 2025	Tier 2	Jan 2025 (anchor)	Around 40% of job skills changing by 2030; 63% of companies planning to reskill rather than hire externally.
WEF, AI and Entry-Level Jobs briefing	Tier 2	Jan 2026	Sector-by-sector field evidence on AI's effect on the on-ramp into knowledge work.
NACE, Job Outlook 2026 Spring Update	Tier 2	Apr 2026	Class of 2026 hiring forecast revised upward as employers rebalance toward using AI to lift entry-level productivity; headline 2026 outlook still weak.

Source	Tier	Date	Key claim used
Brynjolfsson, Li and Raymond, Generative AI at Work	Tier 2	2023 (anchor)	Field experiment, ~5,200 support agents: AI assistance raises resolution rate 15% on average, up to 35% for the least experienced.
OECD, Employment Outlook 2023	Tier 1	Jul 2023 (anchor)	Highest-automation-risk occupations are 27% of OECD employment; high-AI-exposure vacancies skew toward management and business-process skills.
CNBC and Indeed Hiring Lab, AI-skill mentions nearly doubled	Tier 3	Apr 2026	4.2% of full-time early-career postings mention AI skills, nearly double a year earlier.

Cluster 5: The capex reckoning

Source	Tier	Date	Key claim used
Goldman Sachs Global Institute, Tracking Trillions	Tier 2	May 2026	\$7.6tn cumulative AI capex 2026 to 2031, \$765bn in 2026 rising to \$1.6tn in 2031; silicon useful life as the single most influential variable.
Fortune, Big Tech to spend \$700bn on AI in 2026	Tier 3	Apr 2026	Four hyperscalers' combined AI infrastructure capex around \$700bn in 2026.
CNBC, tech AI spending and cash flow	Tier 3	Feb 2026	Amazon 2026 FCF deficit projected \$17bn to \$28bn; Meta FCF down around 90%; Google Cloud Q1 revenue +63%; Microsoft AI annualised \$37bn.
Yahoo Finance and S&P Capital IQ, hyperscalers hit \$700bn	Tier 3	May 2026	Confirms the aggregate \$700bn-plus 2026 hyperscaler capex envelope and the free-cash-flow squeeze.
IMF, Global Economic and Financial Implications of AI	Tier 1	Apr 2026	Scenario lens warning that the gap between AI-prepared and AI-unprepared economies is widening, with fiscal and financial-stability consequences.

Source	Tier	Date	Key claim used
IMF Working Paper 2025/067, AI and Productivity in Europe	Tier 1	2025 (anchor)	Medium-term total-factor-productivity estimates for AI across European economies.

© 2026 Shaping Tomorrow. Strategic Intelligence Report, sample edition. Signal-cluster structural version. The Strategic Intelligence Report format is a Shaping Tomorrow service for senior strategy, foresight, innovation and policy professionals; subscriber cycles include company-specific or sector-specific framing and cycle-over-cycle continuity. See shapingtomorrow.com.