Insurers, Not Model Builders, Will Decide Where AI Agents Can Be Used

AI agents are starting to sign, spend and fail on companies' behalf, and what now limits their use is insurance cover, legal accountability, power supply and running costs rather than what the models can do. Europe's AI rules take full force on 2 August.

Produced using **Shaping Tomorrow's six-stage process for converting external change into decision-ready intelligence**.
Scanned and filtered by AI. **Stress-tested and judged by humans**.

How long you have to get ready

The four actions below, from today to the moment each falls due. Select a pill to see what would change it, or a number to go to the action.

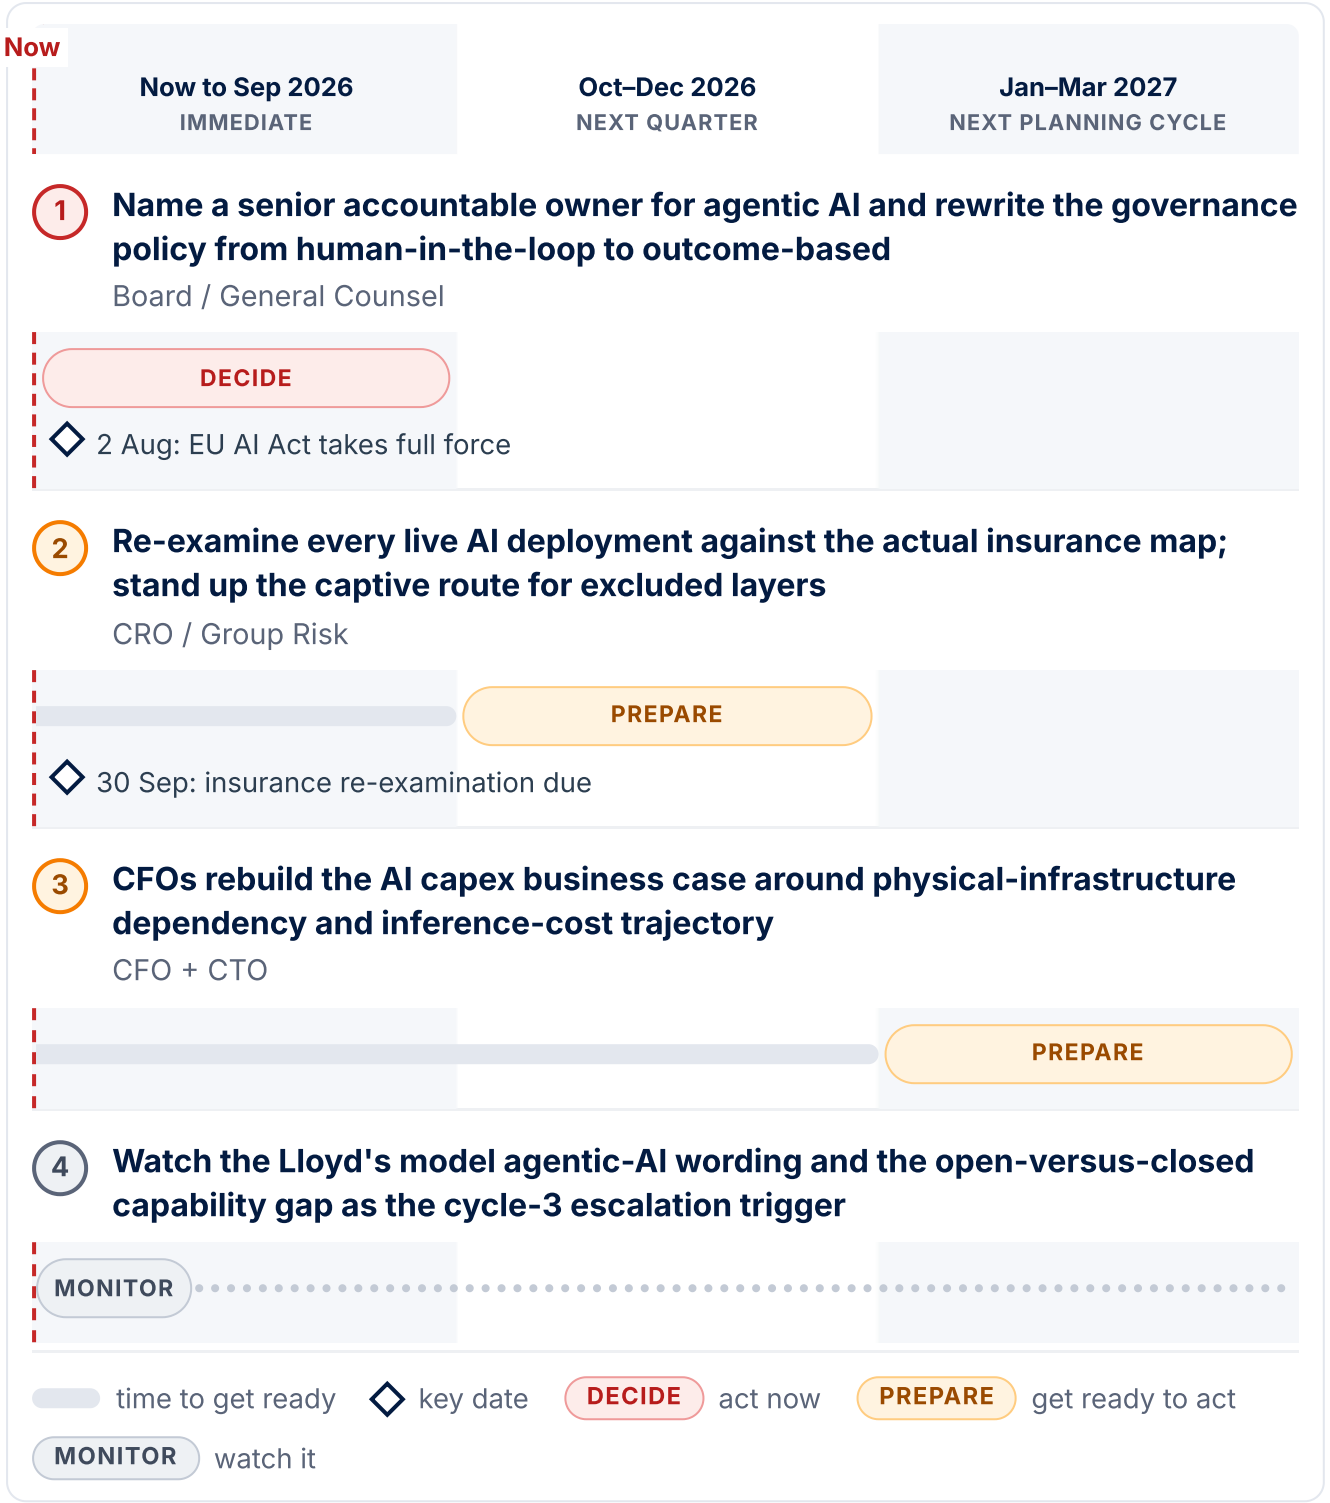
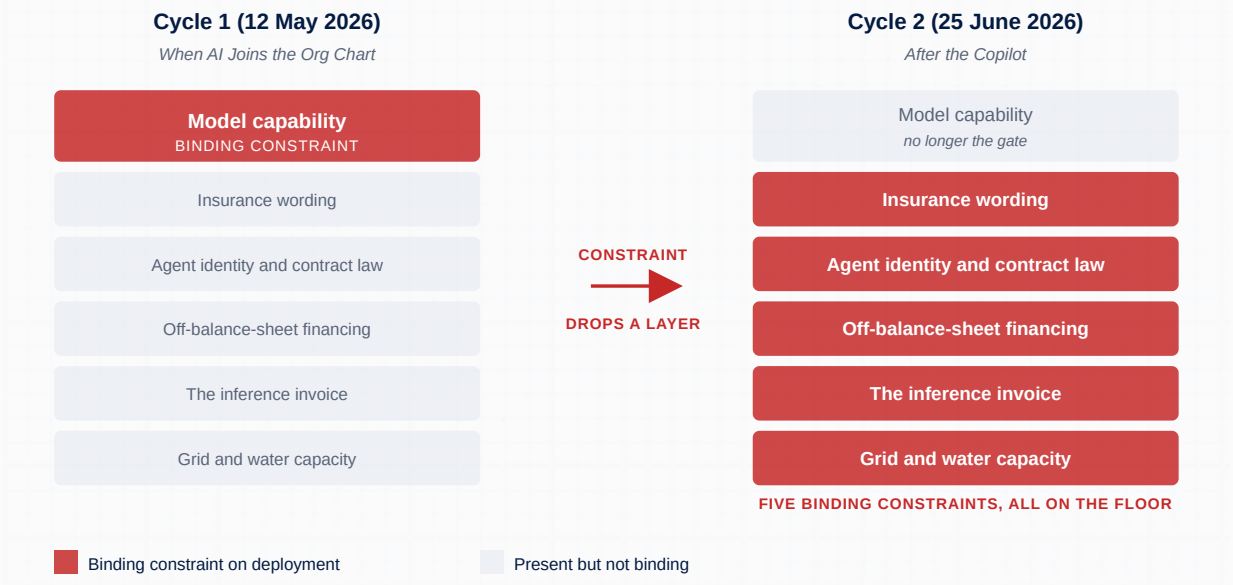


Figure 1: Where the binding constraint on AI deployment actually sits

Cycle 1 looked for the gate at the top of the stack. Cycle 2's evidence located it underneath, in five operational layers no model upgrade can resolve.



Reads across cycle 2 themes 3 (insurance), 1 + 2 (agent identity and oversight), 4 (sovereign capex), 5 (grid and water), and the 13 June Inference Cost Inversion Signal Scan.

Early warnings

Four developments the portfolio caught before consensus formed, in the order they hardened.

LEAD TIME 5 WEEKS AHEAD SIGNAL SCAN, 18 MAY DI CYCLE 2

Agent identity, not model capability, decides which AI is allowed to act on whose behalf

INITIAL SIGNAL

The Signal Scan named the agent-identity-and-authorization stack as the rail that would decide which agents got to act on whose behalf: NIST AI Agent Standards Initiative, the NCCoE concept paper, OpenID Foundation's federal RFI response, Okta and Microsoft Entra agent-ID releases.

NOW EVIDENT

- **Singapore IMDA**: January 2026 Model Governance design treats agent identity as a procurement primitive.
- **OECD**: February 2026 paper.
- **IMF**: April 2026 note on agentic payments.
- **Mastercard**: all-European-issuers Agent Pay enablement, 2 June 2026.

DECIDE

The question is not what an agent can do, it is whether it will be allowed to do it.

LEAD TIME 4 WEEKS AHEAD SIGNAL SCAN, 23 MAY DI CYCLE 2

The binding constraint on agentic AI is insurability, not capability

INITIAL SIGNAL

The Signal Scan flagged the underwriting brake before it was widely read as the deployment gate. US carriers and the standards body for commercial general liability were already inserting AI exclusions into casualty, D&O, E&O and fiduciary lines, while only a thin affirmative-cover market was filling the gap.

NOW EVIDENT

- Roughly **82% of US P&C policies** now carry Verisk/ISO generative-AI exclusion endorsements.
- **Lloyd's Market Association**: AI Adoption Toolkit launched 23 April 2026.
- **Captive practitioners** (Marsh, Hylant, Milliman, Iowa Insurance Division) are absorbing the layers commercial insurers will not write.
- **D&O carriers** are probing board-level AI governance at renewal.

DECIDE For any operator deploying AI in the credit, claims, contracting or trading perimeter.

LEAD TIME 3 WEEKS AHEAD SIGNAL SCAN, 30 MAY DI CYCLE 2

Hyperscaler balance-sheet strength is the wrong place to look for AI financing risk

INITIAL SIGNAL

The Signal Scan caught what BIS Bulletin No 120 (7 January 2026) had quietly named but the wider market had not yet absorbed: a growing share of the AI build-out has migrated off balance sheet into SPVs, private credit and residual-value-guaranteed leases that published debt figures do not reflect.

NOW EVIDENT

- **IMF, FSB, ECB and Bank of England** flagged the concentration between February and May 2026.
- **BIS Bulletin 120**: AI-related investment now exceeds 1% of US GDP.
- **Outstanding private credit to AI firms**: projected at $300bn to $600bn by 2030.

WHY IT MATTERS

A disappointing AI-revenue print has stopped being only a tech-equity correction. It would land on insurers, pension funds and a private-credit book several times the visible bank exposure.

PREPARE For CFOs and asset allocators.

LEAD TIME PRE-CONSENSUS SIGNAL SCAN, 13 JUNE

Cheaper tokens are producing higher AI bills, not lower ones

INITIAL SIGNAL

The Signal Scan surfaced the inference-cost inversion before consensus framing caught up. Per-token prices are collapsing (the "LLMflation" narrative), but agentic and reasoning workloads consume tokens faster than prices fall, so total inference spend on a useful task is rising rather than falling.

THE PATTERN

This is the unit-economics counterpart to the underwriting brake. The consensus narrative is deflation; the enterprise invoice is inflation.

WHY THIS NOW MATTERS

The May briefing's CFO-facing capex stress-test (SI 4) was sized around falling marginal cost. The disconfirming signal is the enterprise invoice.

PREPARE

Trigger to escalate: the first frontier-lab earnings call naming agentic-workload cost inflation explicitly.

Assumption changes since the May briefing

Five explicit cycle-over-cycle assumption changes, in decreasing order of strategic weight.

SHIFT 1 OF 5

"AI is moving from tool to colleague to manager-of-tasks"

PRIOR CYCLE HELD

The cycle's spine, held inside the org-chart frame: AI graduating from tool through colleague to manager-of-tasks.

THIS CYCLE READS

The agent has moved past joining the org chart and started becoming the counterparty itself. Five months of regulatory and payment-rails activity confirm it:

- IMDA, OECD, IMF, UNCITRAL, BIS building identity, authentication, settlement and contract treatment.
- Visa and Mastercard wiring agent-initiated payments into card rails.

FORWARD THESIS

An organisation can have a perfectly clean agent operating model and still be unprepared for the question of who answers when the agent signs.

SHIFT 2 OF 5

"Verification is now a procurement capability, not a policy artefact"

PRIOR CYCLE HELD

A hard EU AI Act 2 August 2026 deadline, with the trust stack arriving on cue and verification handled as a procurement clause.

THIS CYCLE READS

The procurement framing is confirmed; a layer was missing underneath. Verification is now also an insurance procurement question.

- Roughly **82% of US P&C policies** carry generative-AI exclusion endorsements (Munich Re, March 2026).
- Underwriter due-diligence questionnaires now operate as a parallel compliance regime, in many jurisdictions setting governance standards before any AI-specific statute lands.

FORWARD THESIS

The procurement clause was the right artefact. It now sits inside a policy wording that may decline the cover.

SHIFT 3 OF 5

"Human-in-the-loop is the operating control"

PRIOR CYCLE HELD

Implicit in cycle 1's verification theme and in the cycle's reading of the EU AI Act: a human reviewing the agent's output is the operating control.

THIS CYCLE READS

Refuted across four supervisory venues in two months:

- **Bank of England**: recorded firms saying human-in-the-loop is "challenged by the rise of agentic AI".
- **Fed, OCC, FDIC**: carved generative and agentic AI out of US model-risk guidance, 17 April 2026.
- **IOSCO**: 25 May 2026 Supervisory Toolkit notes member experiments using "AI as a judge".
- **FSB**: 10 June 2026 consultation places board and senior-management accountability at the centre.

FORWARD THESIS

The control model has reset from supervising what an agent does to governing what an agent is allowed to do. Any AI governance policy still anchored on human-in-the-loop describes the past, not the present.

SHIFT 4 OF 5

"The $7.6 trillion AI capex envelope binds at hyperscaler free cash flow"

PRIOR CYCLE HELD

The capex reckoning located the binding constraint at the hyperscaler income statement: FCF deficits, depreciation schedules, silicon useful life.

THIS CYCLE READS

Two evidence streams move the constraint elsewhere. The financing layer has migrated off balance sheet; the physical layer is binding.

- **Financing layer**: SPVs, private credit, residual-value-guaranteed leases held by insurers, pension funds and private-credit funds (30 May scan).
- **NERC**: lifted its ten-year US summer peak forecast by 69% in a single revision.
- **Germany's four TSOs**: rationing 270 GW of grid-connection requests against ~94 GW of forecast 2037 need.
- **50+ US cities**: enacted moratoria on new data-centre construction.

FORWARD THESIS

The substation, not the spreadsheet, is the gate.

SHIFT 5 OF 5

"The AI haves and have-nots gap is a measurement lag"

PRIOR CYCLE HELD

The J-curve reading: a credible counter to the haves-and-have-nots distribution story. Adoption was diffuse, returns would follow.

THIS CYCLE READS

The 20 June Signal Scan sharpens the call. At near-saturation adoption, returns are pooling in roughly the fifth of firms that have rebuilt workflows around the technology. The rest run pilots that never reach the income statement. Adoption has gone wide; the return on it has stayed narrow. The divide is scaling and value capture, not access.

FORWARD THESIS

AI portfolios under-deliver not because the tools fail, but because the bolt-on integration pattern caps the return at the integration boundary.

Actions triggered

Each trigger names why it activates now, why it matters, the recommended actions, the primary decision owner, the decision window, and the conditions under which the recommendation should be reviewed.

DECIDE ACTION 1 DI CYCLE 2

Name a senior accountable owner for agentic AI and rewrite the governance policy from human-in-the-loop to outcome-based

TRIGGER CONDITION

This recommendation has been activated because the EU AI Act enters full force on 2 August 2026 with agentic-AI general-purpose obligations attached to identified persons, model-audit and identity requirements are already codified, and no single named executive currently signs for the deployment population inside most enterprises.

PRIMARY DECISION OWNER **Board / General Counsel**

DECISION WINDOW **ImmediateBefore 2 August 2026** DRIVER **EU AI Act full force**

WHY THIS NOW MATTERS

The Fed, OCC and FDIC have moved accountability onto firms. The FSB has placed it on boards and senior management. The EU AI Act takes full force on 2 August 2026 with its human-oversight rules intact. Standing on a stale policy at the next supervisory engagement costs more than writing the new one now.

RECOMMENDED ACTIONS

- Name the senior accountable owner (distinct from the CIO, reporting to the board).
- Replace the human-in-the-loop language with objectives, guardrails, action logging and outcome thresholds.

REVIEW CONDITIONS

This recommendation should be reviewed if the Commission publishes an implementing decision that changes the accountable-person requirement, if a member-state secures a phased-implementation carve-out for agentic systems, or if a major enterprise publishes a viable outcome-based governance policy that becomes a de-facto standard before the deadline.

PREPARE ACTION 2 SIGNAL SCAN, 23 MAY DI CYCLE 2

Re-examine every live AI deployment against the actual insurance map; stand up the captive route for excluded layers

TRIGGER CONDITION

This recommendation has been activated because Lloyd's has confirmed the exclusion perimeter on agentic-AI risk, three major carriers have followed with parallel wording, and the FY27 renewal window closes before any captive stand-up can offer effective cover for the excluded layers.

PRIMARY DECISION OWNER CRO / Group Risk

DECISION WINDOW Next quarterEnd Q3 2026

BUYING WINDOW Next FY27 insurance cycle

WHY THIS NOW MATTERS

The cycle's underwriting evidence is the binding gate on AI deployment, not a downstream compliance task. The insurable wording is still being written; positioning now sets the cover available later.

RECOMMENDED ACTIONS

- By end Q3 2026, test every live or proposed AI deployment carrying material residual exposure against the current cyber, PI and D&O wording. Flag where the policy no longer covers the loss.
- For operators at captive scale: use the FY27 insurance buying cycle to absorb the excluded layers via difference-in-conditions wrappers.
- For sub-captive operators: choose between accepting the exclusion, switching carrier, or restructuring the deployment to avoid the excluded fact pattern.

REVIEW CONDITIONS

A material AI-related loss that triggers an exclusion dispute in the operator's own book, or in a comparable book in the same sector.

PREPARE ACTION 3 SIGNAL SCAN, 30 MAY SIGNAL SCAN, 13 JUNE DI CYCLE 2

CFOs rebuild the AI capex business case around physical-infrastructure dependency and inference-cost trajectory

TRIGGER CONDITION

This recommendation has been activated because the inference invoice has stopped trending down on the hyperscaler curve, grid-connection queues are pushing physical-infrastructure lead times past 24 months, and the FY27 capex round is being locked before either variable stabilises.

PRIMARY DECISION OWNER **CFO + CTO**

DECISION WINDOW **Next planning cycleFY27 planning round**

HORIZON **Three to seven years**

WHY THIS NOW MATTERS

Treat the May briefing's AI-return stress test as superseded. Two cycle-2 evidence streams move the inputs:

- **Physical infrastructure is the binding constraint.** Grid-connection queues run four to seven years; water-rights pushback is municipal-ban territory; behind-the-meter generation triggers utility-classification risk.
- **Inference-cost inversion.** The per-task unit cost of agentic workloads is rising as token prices fall.

RECOMMENDED ACTIONS

- By the FY27 planning round, every AI capex commitment above materiality carries a **named substation**.
- A **named water source**.
- A **named inference-cost trajectory by workload class**, signed off by the CFO.

REVIEW CONDITIONS

A planned data-centre programme rejected at the grid-connection stage, or an inference-cost overrun above the FY27 capex margin.

MONITOR ACTION 4 SIGNAL SCAN, 23 MAY DI CYCLE 2

Watch the Lloyd's model agentic-AI wording and the open-versus-closed capability gap as the cycle-3 escalation trigger

TRIGGER CONDITION

This recommendation has been activated because the Lloyd's model wording is expected in H2 2026 and will set the market baseline, the open-versus-closed capability gap is closing faster than mid-cycle assumptions, and both markers will decide whether cycle 3's counterparty thesis holds or has to be reframed.

TWO NAMED MARKERS FOR THE NEXT SIX MONTHS

- **Lloyd's Market Association sub-limited agentic-AI cyber wording**. Publication by end-2026 would soften the underwriting brake and reopen capacity, reshaping Strategic Implication 2.
- **Open-vs-closed capability gap (Epoch AI monthly index)**. The gap staying inside six months for three consecutive readings would confirm the post-moat reading of Theme 4 and force vendor renegotiation.

REVIEW CONDITIONS

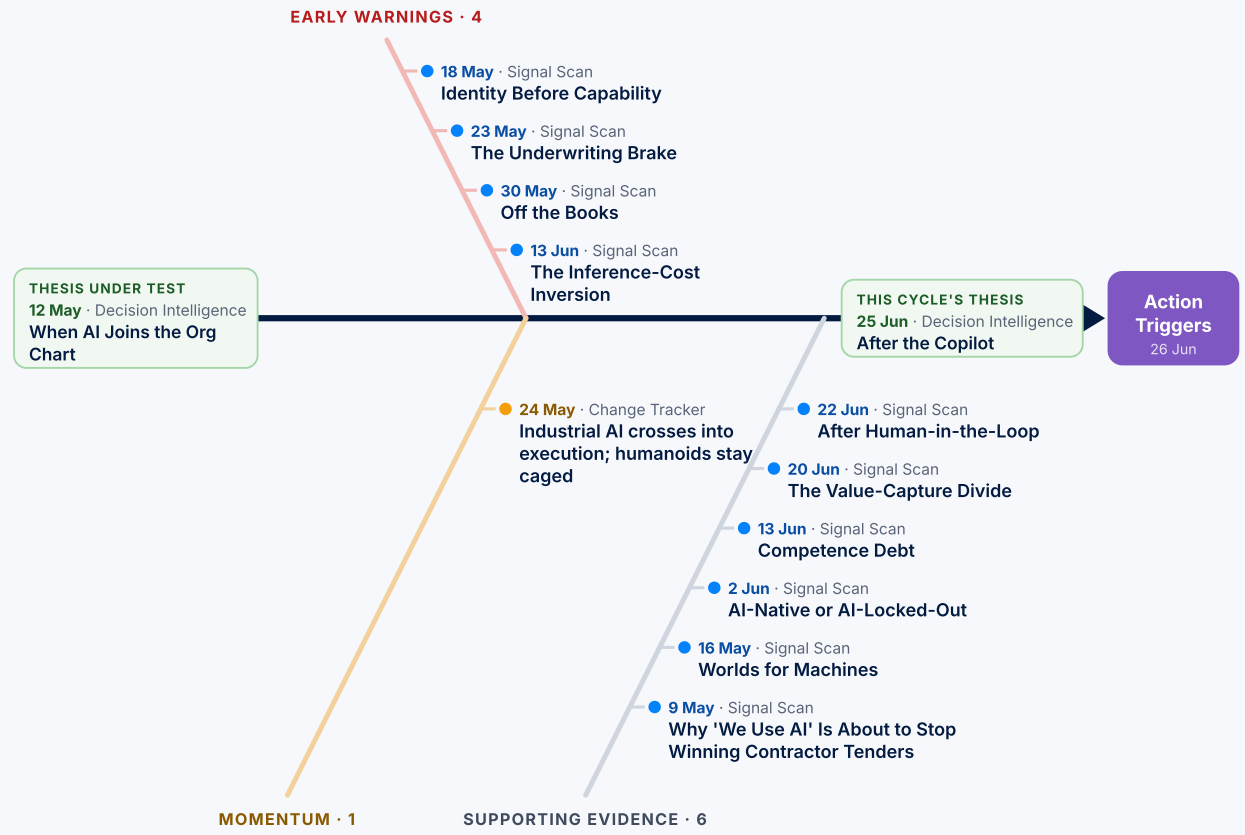
This recommendation should be reviewed if the Lloyd's wording is delayed beyond H2 2026 (buys time), if an open-weight release closes the capability gap decisively (shifts the risk map), or if a major closed-model provider publishes a warranty on its outputs (which would change the counterparty structure the cycle assumed).

> MONITOR Either trigger fires: escalate to Prepare. Absent either, hold posture and revisit at the next Action Triggers cycle.

IF YOU ACT ON ONE THING

Name one senior owner for agentic AI before 2 August, because regulators now hold firms, not systems, accountable for what agents do.

DRAWING ON



Signal Scan Change Tracker Decision Intelligence

Each bone groups the deliverables by what they added, with the oldest at the outer end. Hover over any deliverable for its full title.

HOW THIS WAS MADE

The Decision-Grade Intelligence Method

Shaping Tomorrow's six-stage process for converting external change into decision-ready intelligence.

Scan · Filter · Stress-test · Judge · Trigger · Actionable Intelligence

See full methodology → · How we verify → · Track record →

This Action Triggers brief is the end-of-cycle synthesis of a Shaping Tomorrow intelligence cycle on **AI & Automation**. It reads across the cycle's weekly Signal Scans, monthly Change Trackers and quarterly Decision Intelligence briefing, and surfaces the patterns no single deliverable carries on its own: the early warnings the work caught before consensus formed, the assumptions that have shifted between cycles, and the actions now decision-shaped enough to act on. Continuous intelligence on this topic is available as a Single-Topic or Multi-Topic programme.

Discuss a Tailored Programme

Prepared by Shaping Tomorrow: 26 June 2026