ACTION TRIGGERS · AI & AUTOMATION · 21 AUGUST 2026

SHAPING TOMORROW
DECISION INTELLIGENCE

AI Spending Is Not Earning Its Keep.

Over this cycle, the question changed from who governs AI to who can prove it works. Eight weeks ago the central issue was the floor underneath deployment: insurance wording, agent identity, contract law, the grid queue and the inference invoice. That floor is now occupied. Four financial authorities pulled AI inside a supervisory perimeter in a single month, states passed 28 data-centre laws, and the transparency duties carry fines. What did not arrive is the measurement. Capital spending accelerated past a trillion dollars while the instrument for judging whether any of it pays was withdrawn by the researchers who built it. Below: the signals that hardened first, the five beliefs this cycle changed, and the five decisions now shaped enough to act on.

STRATEGIC PICTURE AT A GLANCE

5 Weak signals became measurable	**5** Strategic assumptions changed	**5** Actions now decision-shaped	**1** Strategic constraint replaced

DRAWING ON

- Decision Intelligence, 20 August 2026 · 'Why AI Projects Die on Value'
- Signal Scan, 15 August 2026 · 'The Memory Queue'
- Signal Scan, 8 August 2026 · 'Steering by Drifting Instruments'
- Change Tracker, 6 August 2026 · 'Supervisors arrive as capex climbs'
- Signal Scan, 1 August 2026 · 'Rationed Discovery'
- Signal Scan, 25 July 2026 · 'Compute Gets a Forward Curve'
- Signal Scan, 18 July 2026 · 'Breadth Without Depth'
- Signal Scan, 11 July 2026 · 'The GPU Depreciation Question'

- Signal Scan, 4 July 2026 · 'The Consent Constraint'
- Change Tracker, 2 July 2026 · 'Capital repricing arrives'
- Signal Scan, 27 June 2026 · 'The Upstream Capture'
- Decision Intelligence, 25 June 2026 · 'After the Copilot'

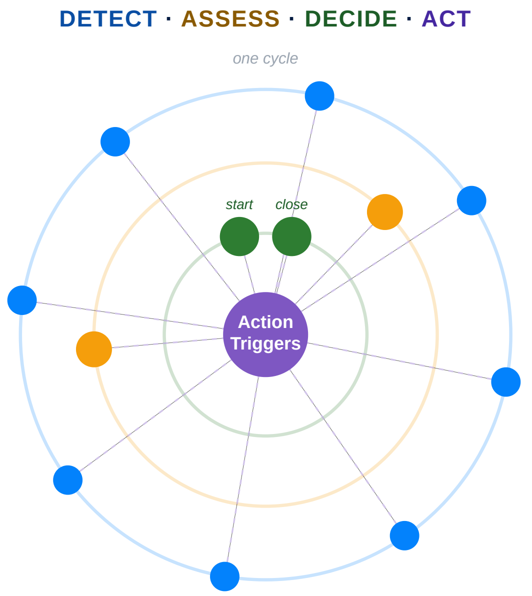

Eight weeks of continuous monitoring, resolved into one decision-grade synthesis at the centre.

Early warnings

Signals the portfolio surfaced before consensus formed, in the order they hardened.

LEAD TIME 1 WEEK AHEAD | SIGNAL SCAN | SIGNAL SCAN

Korea raised interest rates on the strength of an AI export boom

INITIAL SIGNAL

We flagged that AI capital spending had grown large enough to distort the national accounts of the economies supplying it, and that a terms-of-trade windfall was being read as underlying strength.

NOW EVIDENT

- **Concentration**: five Korean firms took 43% of first-quarter export earnings, up from 27%.
- **Forecast error**: the top four AI-hardware net exporters beat first-quarter forecasts by 4.4 percentage points, against a 0.3 point shortfall everywhere else.
- **Policy response**: the Bank of Korea raised its base rate from 2.50 to 2.75% on 16 July, citing export-led growth.
- **Microfoundation**: one supplier posted a 76% operating margin on memory sold into AI demand.

RATE DECISION **16 July 2026** EXPORT CONCENTRATION **43% from five firms**

MONITOR Treasury and planning functions should discount headline growth in AI-hardware exporters before using it as a demand signal.

LEAD TIME 2 WEEKS AHEAD SIGNAL SCAN CHANGE TRACKER

Four financial regulators took charge of AI risk in a single month

INITIAL SIGNAL

We identified that reliance on a handful of cloud, data and model providers had become a system-wide vulnerability, and read the supervisory interest as arriving with the contracts rather than after them.

NOW EVIDENT

Four financial authorities converted AI from market commentary into supervised risk within a single month: the Bank of England Financial Policy Committee and the European Systemic Risk Board on 7 July, the Financial Stability Board's consultation, and ECB Banking Supervision. Frontier-model cyber offence entered the register at the same time, with a named October supervisory deadline.

AUTHORITIES **Four in one month** DEADLINE **31 October 2026**

PREPARE Risk and board functions should assume an AI exposure map will be requested rather than volunteered.

LEAD TIME 3 WEEKS AHEAD SIGNAL SCAN DI CYCLE 3

The gap between AI adoption surveys became the cycle's main finding

INITIAL SIGNAL

We flagged that the three best United States measures of AI adoption diverge four-fold, so there is no single adoption rate: about 18% of firms, 41% of workers, 78% employment-weighted. The early read was that AI is broad but shallow.

NOW EVIDENT

- **Sectoral resolution**: business AI use sat at 19.8% nationally in early May against 39.7% in Information, so national series flatter a narrow base.
- **Depth measured**: 57% of adopting firms use AI in three or fewer business functions.
- **Carried into the briefing**: breadth without depth became two of the five themes in the closing Decision Intelligence cycle, five weeks after we first put the divergence on the page.

DECIDE Boards should stop using a national adoption rate as evidence of competitive position.

LEAD TIME 4 WEEKS AHEAD | SIGNAL SCAN | CHANGE TRACKER

AI spending overtook the cash coming in, and company results confirmed it

INITIAL SIGNAL

We identified that AI capital spending had moved ahead of hyperscaler cash flow, pushing the buildout toward debt and toward circular financing among chipmakers, laboratories and clouds.

NOW EVIDENT

- **Scale**: the five largest hyperscalers are on course for more than a trillion dollars across 2025 and 2026, ahead of earnings and free cash flow.
- **Unevenness**: one hyperscaler's expenses rose 55% to 42 billion dollars against 28% revenue growth.
- **Financing shift**: roughly 570 billion dollars of AI-related debt issuance is projected for 2026, moving toward private credit.

PREPARE Finance functions should treat the funding mix, not the spending total, as the variable that moves first.

LEAD TIME 5 WEEKS AHEAD | SIGNAL SCAN | CHANGE TRACKER

Local opposition became the main obstacle to data centres, and state laws multiplied

INITIAL SIGNAL

We flagged that local political consent had overtaken capital, chips and grid engineering as the binding constraint on the United States buildout, on five enacted state laws and a first statewide moratorium.

NOW EVIDENT

- **Count**: 28 state data-centre laws enacted by 1 July, covering ratepayer protection, audits and tax-break rollbacks, plus two moratorium bills.
- **Register entry**: restrictive siting politics entered the momentum register as a new accelerating signal, having carried no siting signal at all five weeks earlier.
- **Reading**: the constraint is arriving through state energy law rather than through AI law, and it is the first year both parties legislated restrictions.

STATUTES **5 to 28 in five weeks** OPPOSITION **71% oppose a local site**

DECIDE Siting programmes should be resourced as political campaigns, not as permitting exercises.

Assumptions that have shifted

Cycle-over-cycle assumption changes, in decreasing order of strategic weight.

SHIFT 1 OF 5 SPLIT

The same short list of practical obstacles decides whether AI gets deployed

PRIOR CYCLE HELD

Prior thinking placed a single floor beneath AI deployment, made of insurance wording, agent identity, contract law, the grid queue and the inference invoice, and held in its own words that "capability is the wrong place to look for the gate; the floor underneath sets it".

THIS CYCLE READS

There is no single floor. Four different binding constraints were argued within seven weeks, each convincingly and each in a different domain.

- **Local consent**: the constraint has moved from the engineering department to the county board.
- **Grid and chips**: willingness to spend is no longer what binds; turbines, transformers and connections are.
- **An accounting estimate**: the useful life assigned to GPUs, not demand, is the nearer fault line.
- **Contracted memory**: through 2027 the planning variable is secured allocation rather than forecast cost.

FORWARD THESIS

Stop asking what the constraint is and start asking which constraint binds this asset, in this jurisdiction, this year. **The tripwire is any single quarter in which two of the four constraints bind the same programme at once.**

SHIFT 2 OF 5 REPLACED

Regulators are easing off just as investors start repricing AI

PRIOR CYCLE HELD

The prior cycle read the period as "capital repricing arrives as the rulebook stands down": the brake easing on both sides of the Atlantic, with simplification adopted in Europe, one state law stayed and rewritten, and federal pre-emption pressure building in the United States.

THIS CYCLE READS

Light touch above, hard law below. The easing happened at the level that attracts headlines and reversed at the level that binds operations.

- **No pre-emption**: no federal statute passed, while 29 states legislated anyway, 109 AI laws in total.
- **Duties live**: transparency obligations applied from 2 August, with penalties reaching 15 million euros or 3% of turnover.
- **Near date**: a grace period closes on 2 December 2026 for content already on the market, which makes this an autumn procurement question rather than a 2027 one.

FORWARD THESIS

Read deregulation announcements as signals about attention, not about obligation. **The tripwire is the first enforcement action taken under the transparency duties.**

SHIFT 3 OF 5 SPLIT

AI spending is not earning its keep, and the problem is how companies adopt it

PRIOR CYCLE HELD

Prior thinking held that enterprise AI was failing to convert spending into measured value, and that, in the phrase carried at the time, "the gap is organisational, not technological".

THIS CYCLE READS

The reading splits by altitude, and the split is the most useful thing in this section.

- **At sector level, overtaken**: the first broad measured productivity comparison found 34% growth in the most AI-exposed sectors against 24% in the least, with a 62% wage premium on AI skills.
- **At deployment level, hardened**: the field's most rigorous experiment was withdrawn by its own authors, and self-report overstates the time effect by about 40 percentage points.

FORWARD THESIS

Sector-level gains and deployment-level blindness can both be true, and an organisation cannot borrow the first to justify the second. **The tripwire is any credible replacement for the withdrawn experimental design.**

SHIFT 4 OF 5 OVERTURNED

The late-June selloff marked the top of the AI spending boom

PRIOR CYCLE HELD

A late-June selloff, with the index off more than 2% and one memory maker off more than 13, was read at the time as the moment "the one-way climb cracked".

THIS CYCLE READS

The inflection did not hold. Spending kept climbing and the bill landed unevenly, which moved capital expenditure back to accelerating at the highest heat on the register within five weeks.

FORWARD THESIS

A price move in AI-exposed equities is not evidence about AI capital plans, which are set on multi-year supply contracts. **The tripwire is a hyperscaler guiding capital expenditure down rather than a quarter of falling share prices.**

SHIFT 5 OF 5 REPLACED

Free open-weight models will undercut the big AI labs on price

PRIOR CYCLE HELD

Open weights were carried as a disruptive opportunity under the heading "open-weight commoditisation shock", on the reading that near-zero inference cost would compete pricing power away from the closed laboratories.

THIS CYCLE READS

The same signal now sits on the register as a disruptive risk, after an export-control intervention forced two frontier open-weight models fully offline. The competitive question turned into a control question.

FORWARD THESIS

Treat open-weight availability as a jurisdictional variable rather than a market one, and hold a fallback for any dependency on a specific open model. **The tripwire is a second export-control action against an open-weight release.**

Actions triggered

Each trigger names why it activates now, why it matters, the recommended actions, the primary decision owner, the decision window, and the conditions under which the recommendation should be reviewed.

DECIDE ACTION 1 SIGNAL SCAN

Secure memory allocation for every 2027 hardware programme before the contracting window closes

TRIGGER CONDITION

This recommendation has been activated because three suppliers have contracted years of DRAM output to AI customers, because contract prices have moved by a multiple rather than a margin and one supplier is posting a 76% operating margin, and because the resulting shortage is being settled by allocation rather than cleared by price, which means a purchase order is no longer a claim on supply.

PRIMARY DECISION OWNER **Chief Operating Officer, with Procurement**

DECISION WINDOW IMMEDIATE **Before 2027 build slots are committed**

WHY THIS NOW MATTERS

Allocation markets punish the buyer who plans on price. Any product with memory in it, from vehicles to security hardware to handsets, is now competing for capacity against a customer with better margins and longer contracts.

RECOMMENDED ACTIONS

- **Immediate:** list every product line with a memory dependency and establish whether its 2027 supply is contracted or forecast.
- **Immediate:** convert forecast positions into contracted ones where the volume justifies it, accepting a worse price for a secured quantity.
- **Follow-on:** qualify a second source on the components where automotive-grade or long-life qualification is the binding delay rather than the part itself.

REVIEW CONDITIONS

This recommendation should be reviewed if suppliers add capacity that reaches the market before mid-2027, if AI demand for high-bandwidth memory cools enough to release standard capacity, or if your own product roadmap moves to designs with materially lower memory content.

PREPARE ACTION 2 SIGNAL SCAN CHANGE TRACKER

Produce the AI exposure map before a supervisor asks for it

TRIGGER CONDITION

This recommendation has been activated because four financial authorities brought AI inside a formal supervisory perimeter within a single month, because a named supervisory deadline now falls at the end of October, and because the concentration being supervised is third-party dependency on a small number of cloud, data and model providers, which is a map most organisations have never drawn.

PRIMARY DECISION OWNER Chief Risk Officer, reporting to the Board

DECISION WINDOW NEXT QUARTER Ahead of the 31 October supervisory deadline

WHY THIS NOW MATTERS

Supervisory attention has moved from what a model does to whom the firm depends on. An exposure map assembled under deadline is worse than one assembled deliberately, and the first request will not come with notice.

RECOMMENDED ACTIONS

- **Immediate:** map every material AI dependency to its underlying cloud, model and data provider, including the ones reached through a vendor rather than directly.
- **Next quarter:** identify the single points of failure the map exposes and name the substitution path for each.
- **Follow-on:** place the map on a refresh cycle rather than treating it as a one-off exercise, since the dependencies move with vendor architecture.

REVIEW CONDITIONS

This recommendation should be reviewed if a supervisor publishes a template that supersedes an internally designed map, if your sector is explicitly scoped out of the perimeter, or if a concentration threshold is set low enough that the exercise becomes a compliance return rather than a risk instrument.

PREPARE ACTION 3 SIGNAL SCAN SIGNAL SCAN

Test the AI capital plan against a public compute price and a shorter hardware life

TRIGGER CONDITION

This recommendation has been activated because reported profitability across the buildout rests on a discretionary four-to-six-year useful life for hardware that turns over economically in two to three, because exchange-listed compute futures will put a public forward curve on residual-value assumptions currently held inside depreciation schedules and loan books, and because the financing mix has shifted toward private credit at a projected 570 billion dollars for the year.

PRIMARY DECISION OWNER **Chief Financial Officer**

DECISION WINDOW NEXT PLANNING CYCLE **Into the 2027 budget round**

WHY THIS NOW MATTERS

A published curve converts a private estimate into a mark. The exposure is not only for owners of the hardware: it reaches anyone whose supplier, landlord or counterparty is carrying the same assumption.

RECOMMENDED ACTIONS

- **Next quarter:** run the capital plan at a three-year hardware life and record what breaks.
- **Next quarter:** identify contracts whose pricing assumes compute costs falling, and establish what happens if the curve prints flat or upward.
- **Follow-on:** ask key AI suppliers what useful life sits behind their pricing, and treat an evasive answer as information.

REVIEW CONDITIONS

This recommendation should be reviewed if the futures market fails to establish liquidity, if hardware rental prices keep rising, which is what the evidence did this cycle and which cuts against the depreciation-reset case, or if an accounting standard-setter removes the discretion and settles the question centrally.

PREPARE ACTION 4 SIGNAL SCAN CHANGE TRACKER

Treat data-centre siting as a political campaign, not a permitting exercise

TRIGGER CONDITION

This recommendation has been activated because enacted state data-centre laws multiplied from five to 28 in five weeks, because those statutes arrive through energy law rather than AI law and therefore land on ratepayer protection, audits and tax-break rollbacks, and because it is the first year in which both parties legislated restrictions, which removes the assumption that an election changes the direction.

PRIMARY DECISION OWNER **Head of Government Affairs, with Development**

DECISION WINDOW NEXT QUARTER **Ahead of the 2027 state legislative sessions**

WHY THIS NOW MATTERS

Roughly 130 billion dollars of projects were blocked or delayed in a single quarter, which is a larger number than most siting teams have ever had to argue against. Consent is now a schedule input with a cost attached.

RECOMMENDED ACTIONS

- **Immediate:** re-baseline every site in the pipeline against the enacted statute in its own state rather than against a national assumption.
- **Next quarter:** put a cost-allocation position on paper before a regulator writes one, since the ratepayer question is the one that carries the politics.
- **Follow-on:** build the flexibility case, because the same load being legislated against can be offered back to the grid as a service.

REVIEW CONDITIONS

This recommendation should be reviewed if federal pre-emption passes as statute rather than as executive direction, if a state that legislated restrictions reverses under fiscal pressure, or if grid operators begin contracting data-centre flexibility at a price that changes the local argument.

MONITOR ACTION 5 SIGNAL SCAN

Keep watch on who controls the parts and standards behind physical AI

TRIGGER CONDITION

This recommendation has been activated because the contest in physical AI is being settled upstream in training data, components and standards rather than in finished robots, because one state has pre-financed all three layers years before a volume market exists, including a national standard system built with more than 120 institutions and roughly 90% of permanent-magnet processing, and because this subject appears once in the cycle and never enters the momentum register at all.

PRIMARY DECISION OWNER **Head of Strategy**

DECISION WINDOW MONITOR ONLY **Reassess at the next cycle**

WHY THIS NOW MATTERS

Component concentration decides cost before a market exists: building one humanoid programme without those suppliers has been costed at roughly three times the price. That is a supply-chain question for anyone buying automation, not only for robot makers.

RECOMMENDED ACTIONS

- **Immediate:** establish whether any automation on your 2027 roadmap depends on actuators, sensors or magnets from a single jurisdiction.
- **Follow-on:** track the standards work rather than the product launches, because the standard is what sets the switching cost.

REVIEW CONDITIONS

Reinstate this as Prepare if the 10,000-unit deployment target for the end of 2026 is met on schedule, if physical AI enters the momentum register as an accelerating signal, or if an export control is placed on actuator or magnet supply. Weaken it if humanoid deployment slips materially and the component question loses its urgency.

The body of work, in one sentence

Over this cycle, the constraint on AI moved from permission to proof. The previous question was whether the rules, the insurers and the grid would let deployment proceed; supervisors have now answered it, and the answer is a perimeter rather than a brake. The new question is whether anyone can show a return, and on that the portfolio surfaced a measurement vacuum rather than a verdict. Plan against the second question, because the first one is settled and the second one is not.

This briefing is not another report. It is the synthesis produced after continuously monitoring signals, tracking change, testing assumptions and updating scenarios throughout the cycle on **AI & Automation**. Its purpose is to identify the points at which accumulated evidence has become decision-changing: the early warnings the work caught before consensus formed, the assumptions that have shifted between cycles, and the actions now decision-shaped enough to act on. Continuous intelligence on this topic is available as a Single-Topic or Multi-Topic programme.

Discuss a Tailored Programme

Prepared by Shaping Tomorrow: 21 August 2026